StockDrop Performance Analysis

Source cohort: 2026-05-10-package

Generated 2026-05-11 17:01

StockDrop Performance Analysis — 2026-02-01 cohort

Generated 2026-05-10 12:35. Cohort: **359 decisions**.

Executive summary

Cohort of **359 decisions** with `decision_date >= 2026-02-01`, evaluated against yfinance OHLC bars cached locally. Headline 4-week metrics on BUY/BUY_LIMIT signals: win rate **100.0%**, avg return **+27.06%** (n=12).

Headlines

- BUY/BUY_LIMIT 4w win rate: **100.0%** (avg +27.06%, n=12)
- BUY_LIMIT fill rate: **56.0%** (28/50 filled), avg 4w on filled: **+15.17%**
- Median days to recover (recovered cohort): **2.0 days** (n=174)

SPY median over the same calendar windows: day-5 +1.61%, day-10 +3.26%, day-20 +7.60%.

Excess vs SPY at day-20 (group median minus SPY median): **ENTER_NOW** +10.58% (n=43), **ENTER_LIMIT** +6.63% (n=42), **AVOID** +0.37% (n=180), **NEUTRAL** -6.05% (n=67).

1. Equity curve

Equal-weight cumulative growth assuming each ENTER_NOW/ENTER_LIMIT decision is held for 4 weeks at its 4w return.

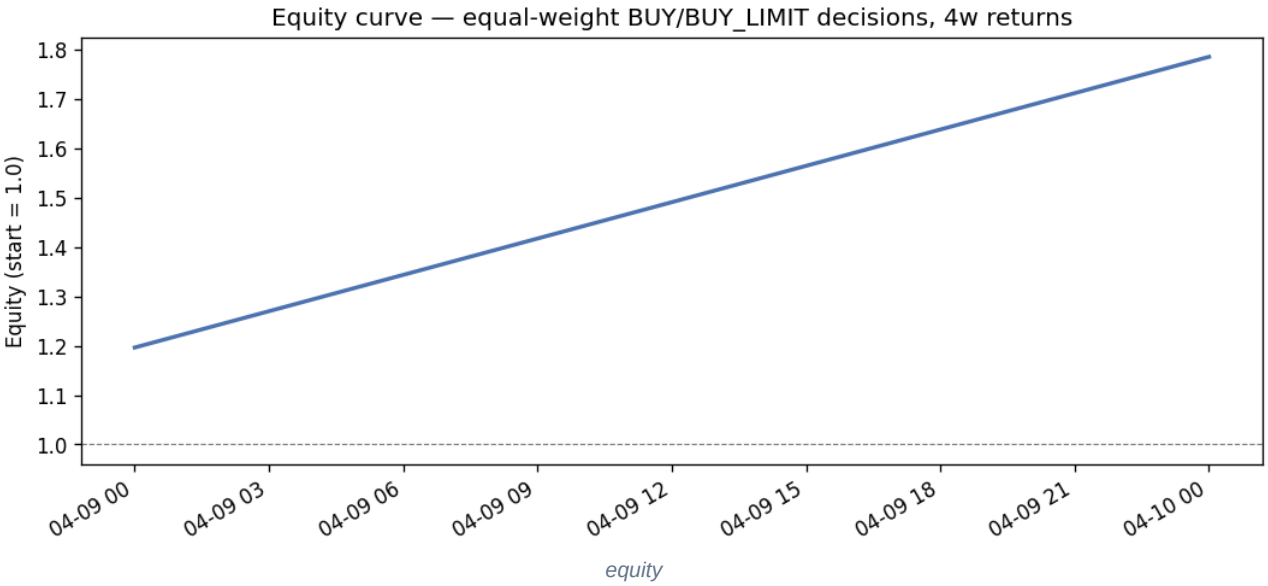

2. Verdict performance

2.1 AI council (PM) intent

intent	count	win_rate	win_rate_se	win_rate_ci_low	win_rate_ci_high	avg_return	avg_return_se	avg_return_ci_low	avg_return_ci_high	median_return	std_return
AVOID	17	0.588	0.119365	0.360054	0.783889	0.109	0.0471709	0.00894561	0.208941	0.08	0.194
ENTER_NOW	8	1	0	0.675592	1	0.221	0.0605294	0.0776654	0.363924	0.182	0.171

intent	count	win_rate	win_rate_se	win_rate_ci_low	win_rate_ci_high	avg_return	avg_return_se	avg_return_ci_low	avg_return_ci_high	median_return	std_return
NEUTRAL	8	0.625	0.171163	0.305742	0.863156	0.138	0.148552	-0.213284	0.489257	0.015	0.42
ENTER_LIMIT	4	1	0	0.510109	1	0.37	0.267987	-0.482719	1.22299	0.142	0.536

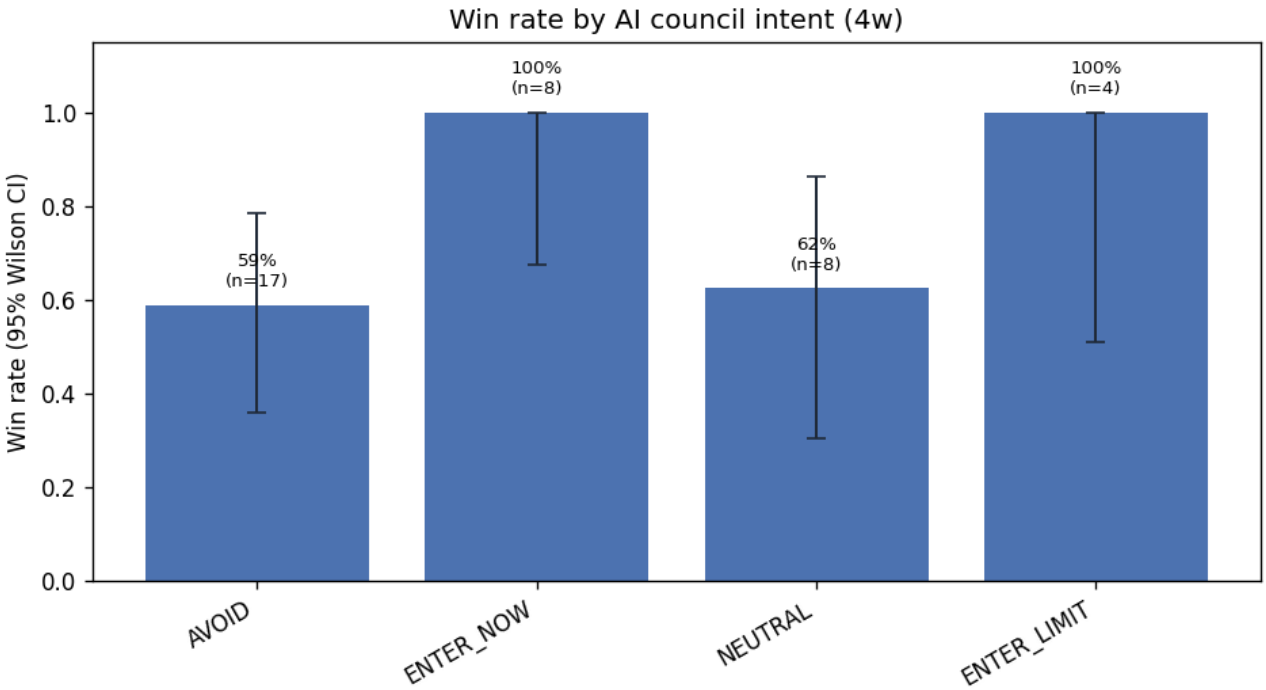

wr_intent

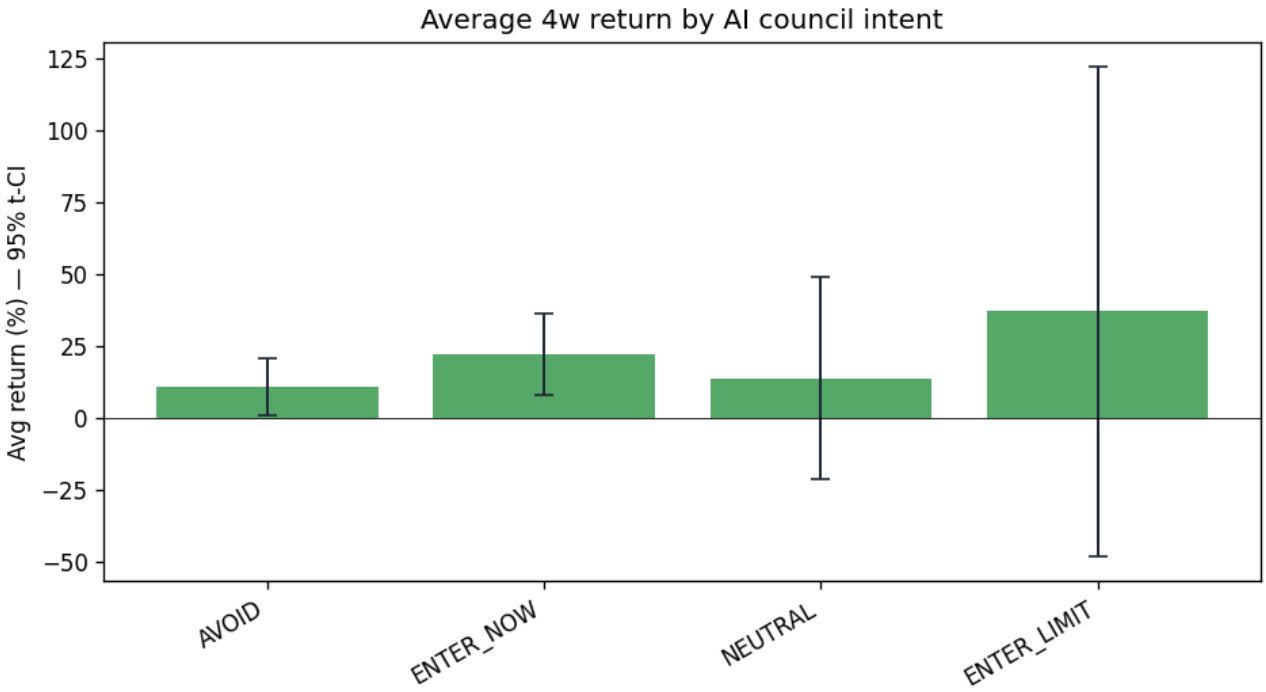

ar_intent

2.2 Deep Research verdict

deep_research_verdict	count	win_rate	win_rate_se	win_rate_ci_low	win_rate_ci_high	avg_return	avg_return_se	avg_return_ci_low	avg_return_ci_high	median_return	std_return
	27	0.667	0.0907218	0.478248	0.813567	0.141	0.0536716	0.0304798	0.251127	0.047	0.279
BUY_LIMIT	6	1	0	0.609666	1	0.335	0.171998	-0.107388	0.776884	0.173	0.421
AVOID	3	1	0	0.438503	1	0.152	0.0383088	-0.012357	0.317302	0.148	0.066
BUY	1	0	0	0	0.793451	-0.065	nan	nan	nan	-0.065	

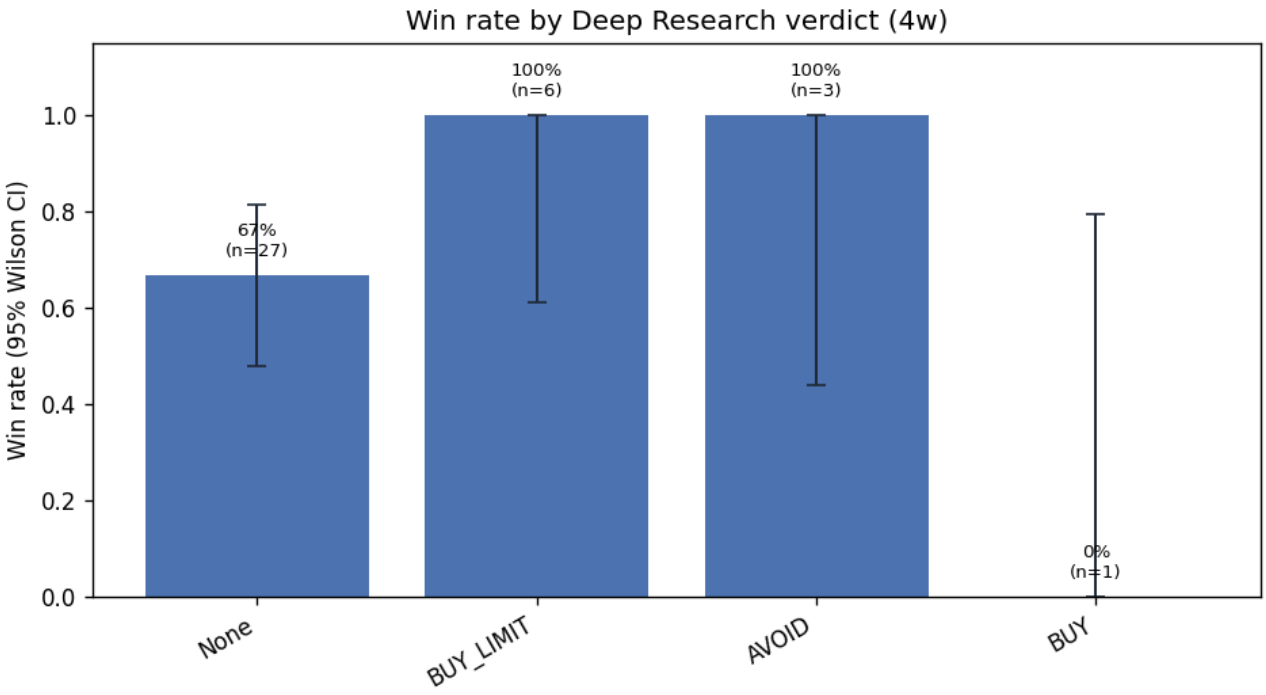

wr_dr

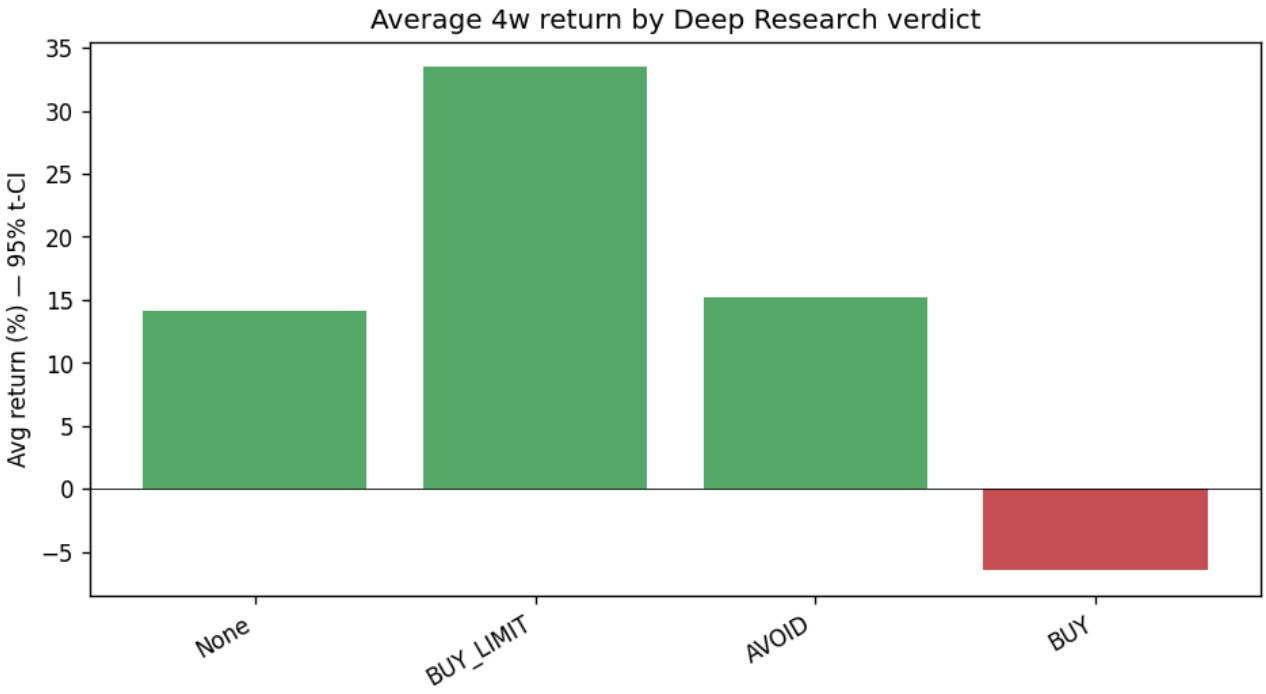

ar_dr

3. Statistical significance

None of the 3 pairwise PM-intent comparisons reach FDR-adjusted p<0.05. Sample sizes are small (n=17, 17, 8 per group), so the apparent gaps in win rate are not yet distinguishable from noise.

3.1 Pairwise tests on AI council intent

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	ENTER_NOW	17	8	-11.19%	-0.6	0.165	0.495	0.097	0.146	—
AVOID	NEUTRAL	17	8	-2.90%	-0.1	0.857	0.857	0.366	0.366	—
ENTER_NOW	NEUTRAL	8	8	+8.28%	0.26	0.618	0.857	0.024	0.072	—

3.2 Pairwise tests on Deep Research verdict

DR verdicts: **none of the 1 pairwise comparisons** are significant — DR groups have n=3–6 each, far below what's needed to detect even large effects.

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	BUY_LIMIT	3	6	-18.23%	-0.51	0.344	0.344	0.905	0.905	—

Interpretation. Welch's t-test compares group means under the (relaxed) assumption that variances may differ. Mann-Whitney U is rank-based and works even when returns are skewed. Both p-values are FDR-adjusted (Benjamini-Hochberg) to control the false-discovery rate across the family of comparisons.

3.3 Pairwise tests at 1-week and 2-week horizons

Pairwise PM intent at 1w (n_per_group is much larger):

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	ENTER_LIMIT	154	27	-0.00%	-0	0.999	0.999	0.903	0.903	—
AVOID	ENTER_NOW	154	32	-1.46%	-0.16	0.291	0.836	0.369	0.903	—
AVOID	NEUTRAL	154	55	+0.58%	0.06	0.669	0.893	0.613	0.903	—
ENTER_LIMIT	ENTER_NOW	27	32	-1.46%	-0.22	0.418	0.836	0.489	0.903	—
ENTER_LIMIT	NEUTRAL	27	55	+0.58%	0.07	0.745	0.893	0.809	0.903	—
ENTER_NOW	NEUTRAL	32	55	+2.04%	0.27	0.199	0.836	0.251	0.903	—

Pairwise PM intent at 2w:

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	ENTER_LIMIT	113	16	-2.62%	-0.24	0.373	0.702	0.604	0.612	—
AVOID	ENTER_NOW	113	16	-4.33%	-0.41	0.027	0.163	0.066	0.28	—
AVOID	NEUTRAL	113	35	-3.09%	-0.29	0.108	0.323	0.093	0.28	—
ENTER_LIMIT	ENTER_NOW	16	16	-1.71%	-0.2	0.585	0.702	0.418	0.612	—
ENTER_LIMIT	NEUTRAL	16	35	-0.47%	-0.05	0.88	0.88	0.612	0.612	—
ENTER_NOW	NEUTRAL	16	35	+1.24%	0.15	0.58	0.702	0.577	0.612	—

Pairwise DR verdict at 1w:

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	BUY	15	6	+1.00%	0.17	0.696	0.696	0.733	0.733	—
AVOID	BUY_LIMIT	15	18	-1.82%	-0.28	0.43	0.645	0.504	0.733	—
BUY	BUY_LIMIT	6	18	-2.83%	-0.44	0.279	0.645	0.415	0.733	—

Pairwise DR verdict at 2w:

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	BUY	7	4	+0.82%	0.13	0.867	0.937	0.788	1	—
AVOID	BUY_LIMIT	7	10	-0.28%	-0.04	0.937	0.937	1	1	—
BUY	BUY_LIMIT	4	10	-1.10%	-0.12	0.839	0.937	0.839	1	—

4. R/R ratio vs realized return

4.1 AI council R/R

bucket	count	win_rate	win_rate_se	win_rate_ci_low	win_rate_ci_high	avg_return	avg_return_se	avg_return_ci_low	avg_return_ci_high	median_return	std_return
<1	23	0.609	0.101764	0.407855	0.778424	0.075	0.0311253	0.0107923	0.139892	0.021	0.149
1-2	6	0.833	0.152145	0.436497	0.969947	0.098	0.0348772	0.00879879	0.188108	0.122	0.085
2-3	5	1.000	0	0.565518	1	0.405	0.213448	-0.187151	0.998104	0.148	0.477
>=3	0		nan	nan	nan		nan	nan	nan		

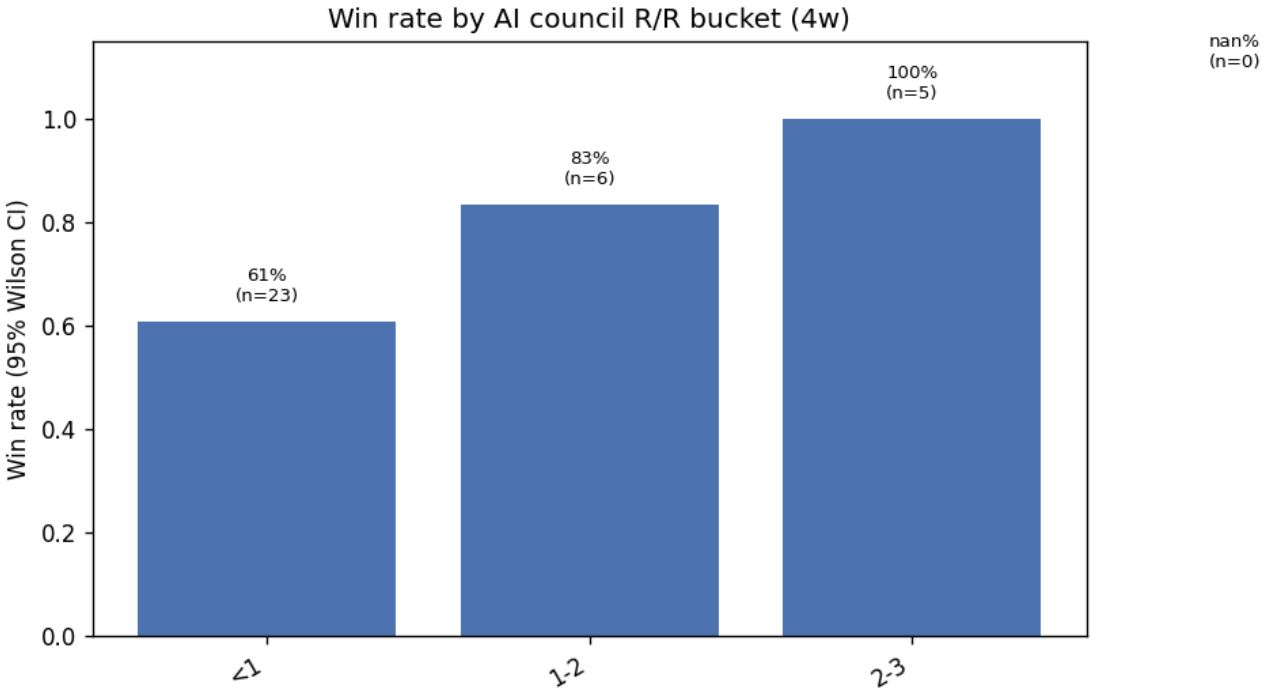

wr_pmrr

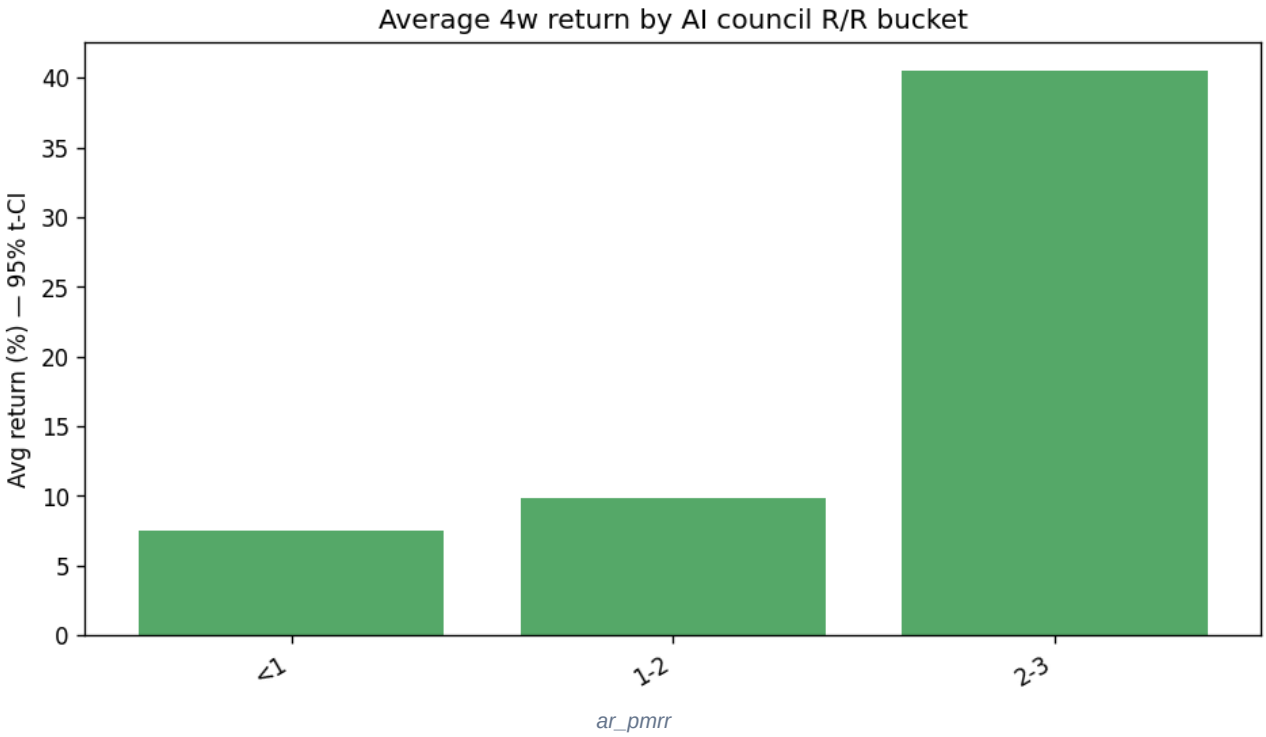

PM R/R vs 4w return (n=34): Pearson r=+0.286 (p=0.101) 95% CI [-0.06, +0.57], Spearman ρ=+0.012 (p=0.945) 95% CI [-0.33, +0.35]. Neither correlation reaches p<0.05.

Correlation at multiple horizons:

horizon	n	Pearson r	Pearson p	Pearson 95% CI	Spearman ρ	Spearman p	Spearman 95% CI
1w	243	-0.005	0.944	[-0.130, +0.121]	0.035	0.592	[-0.092, +0.160]
2w	161	0.077	0.331	[-0.079, +0.229]	0.104	0.188	[-0.052, +0.255]
4w	34	0.286	0.101	[-0.058, +0.569]	0.012	0.945	[-0.327, +0.349]

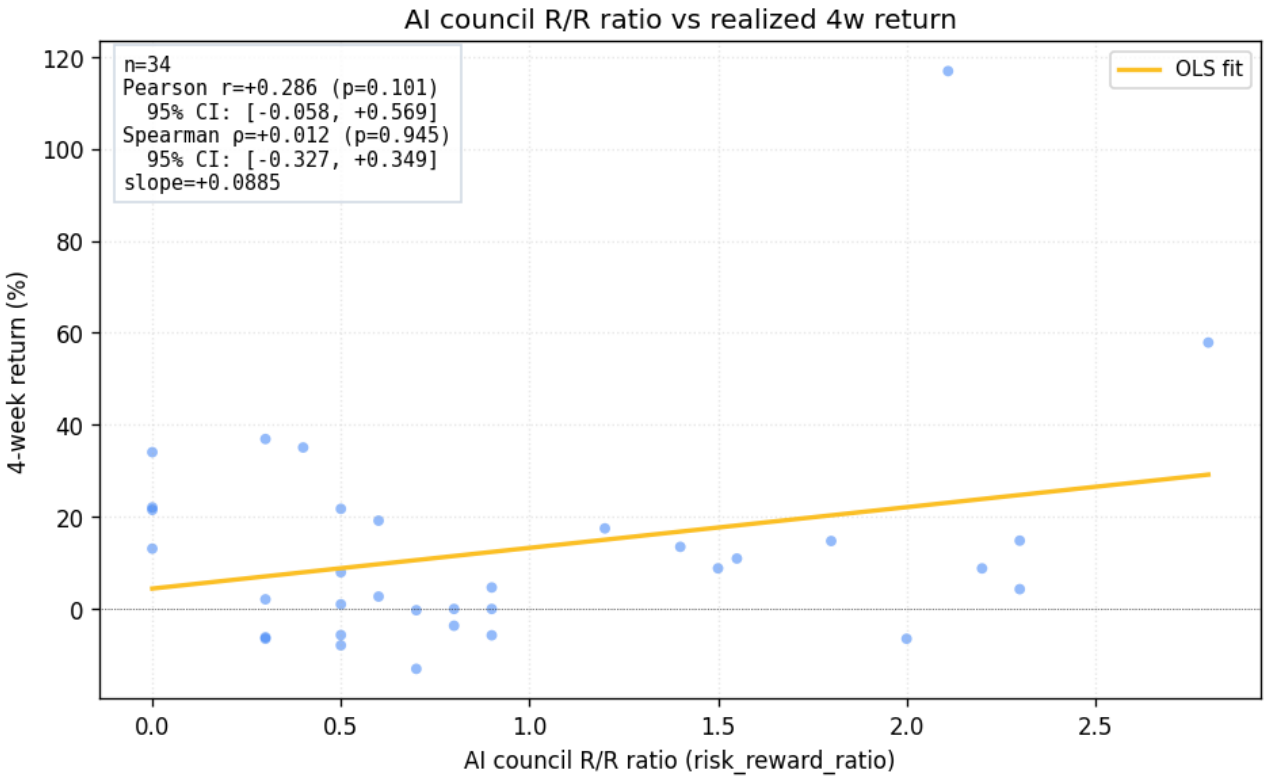

corr_pm

4.2 Deep Research R/R

bucket	count	win_rate	win_rate_se	win_rate_ci_low	win_rate_ci_high	avg_return	avg_return_se	avg_return_ci_low	avg_return_ci_high	median_return	std_return
<1	4	1.000	0	0.510109	1	0.227	0.0431036	0.0897685	0.364119	0.218	0.086
2-3	3	1.000	0	0.438503	1	0.453	0.359045	-1.09141	1.99828	0.148	0.622
1-2	2	0.500	0.353553	0.0945312	0.905469	0.022	0.0871786	-1.0853	1.13011	0.022	0.123
>=3	0		nan	nan	nan		nan	nan	nan		

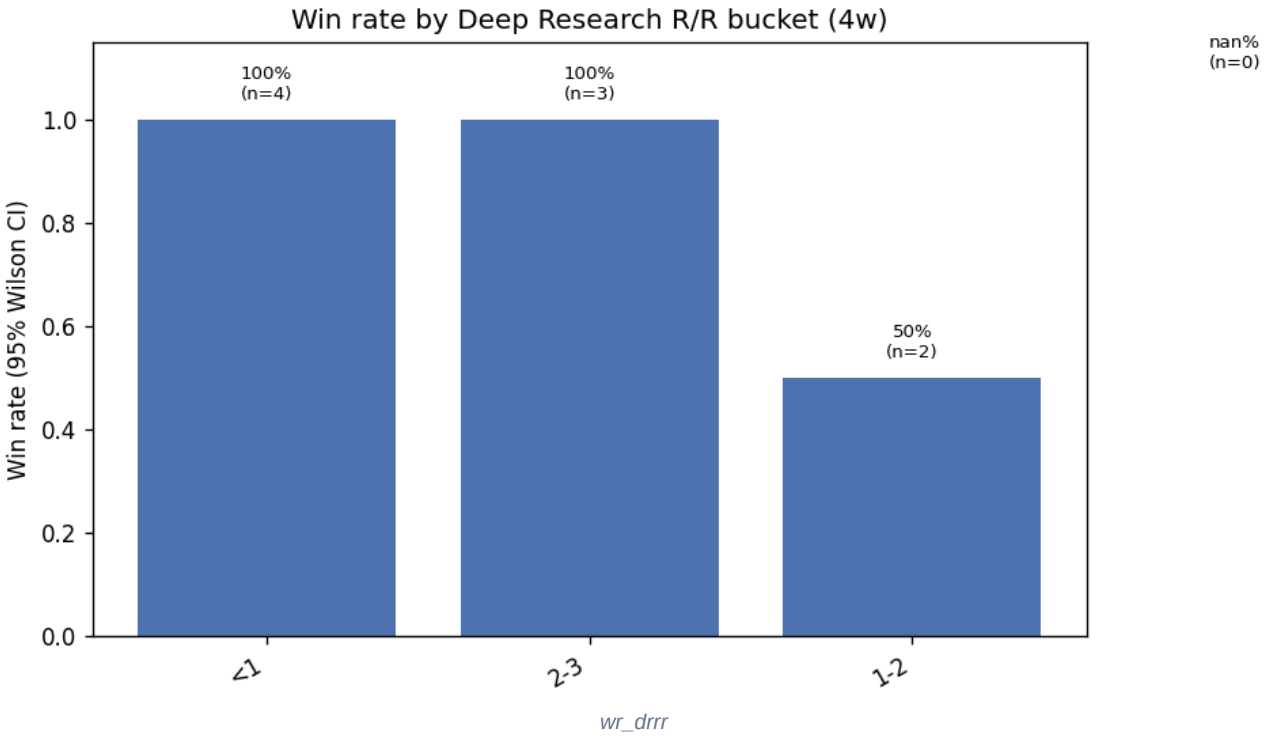

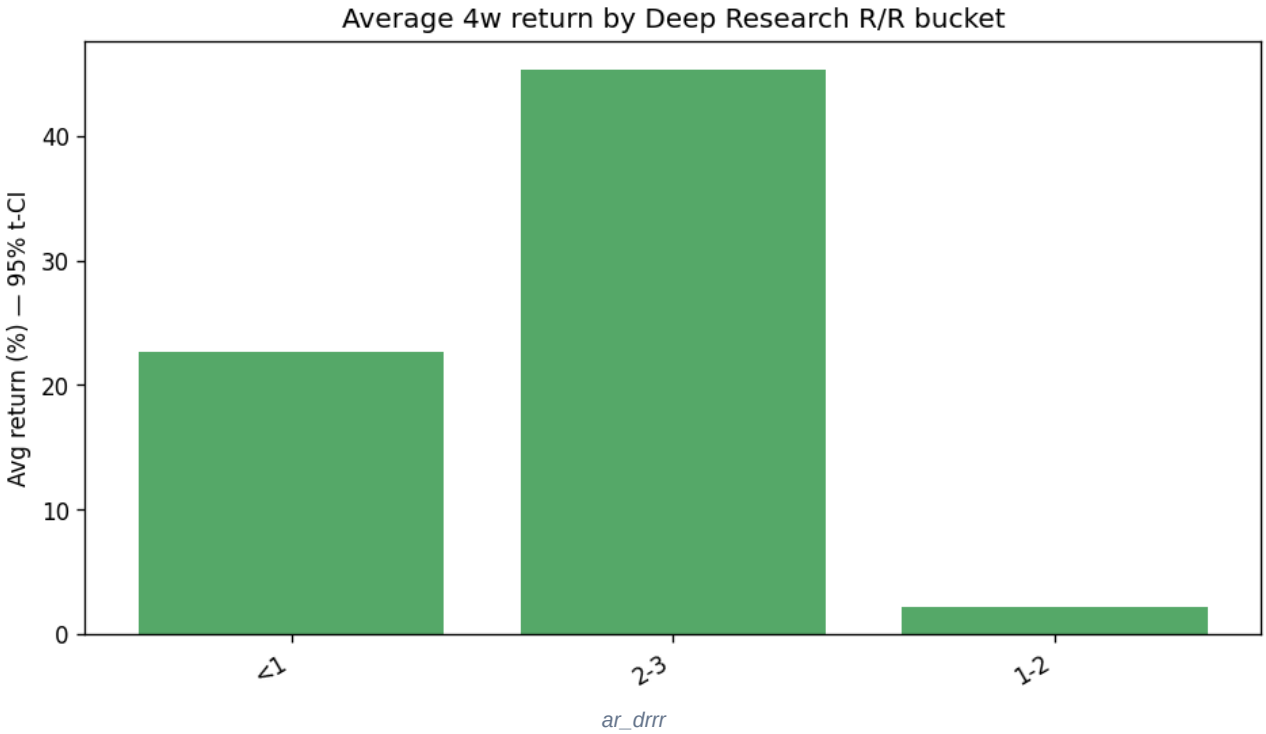

DR R/R vs 4w return (n=9): Pearson r=+0.097 (p=0.803), Spearman ρ=-0.332 (p=0.382). Sample size is small — conclusions are tentative.

Correlation at multiple horizons:

horizon	n	Pearson r	Pearson p	Pearson 95% CI	Spearman ρ	Spearman p	Spearman 95% CI
1w	37	-0.199	0.237	[-0.492, +0.133]	-0.153	0.367	[-0.456, +0.182]

horizon	n	Pearson r	Pearson p	Pearson 95% CI	Spearman ρ	Spearman p	Spearman 95% CI
2w	18	-0.344	0.162	[-0.699, +0.146]	-0.467	0.051	[-0.778, +0.026]
4w	9	0.097	0.803	[-0.606, +0.715]	-0.332	0.382	

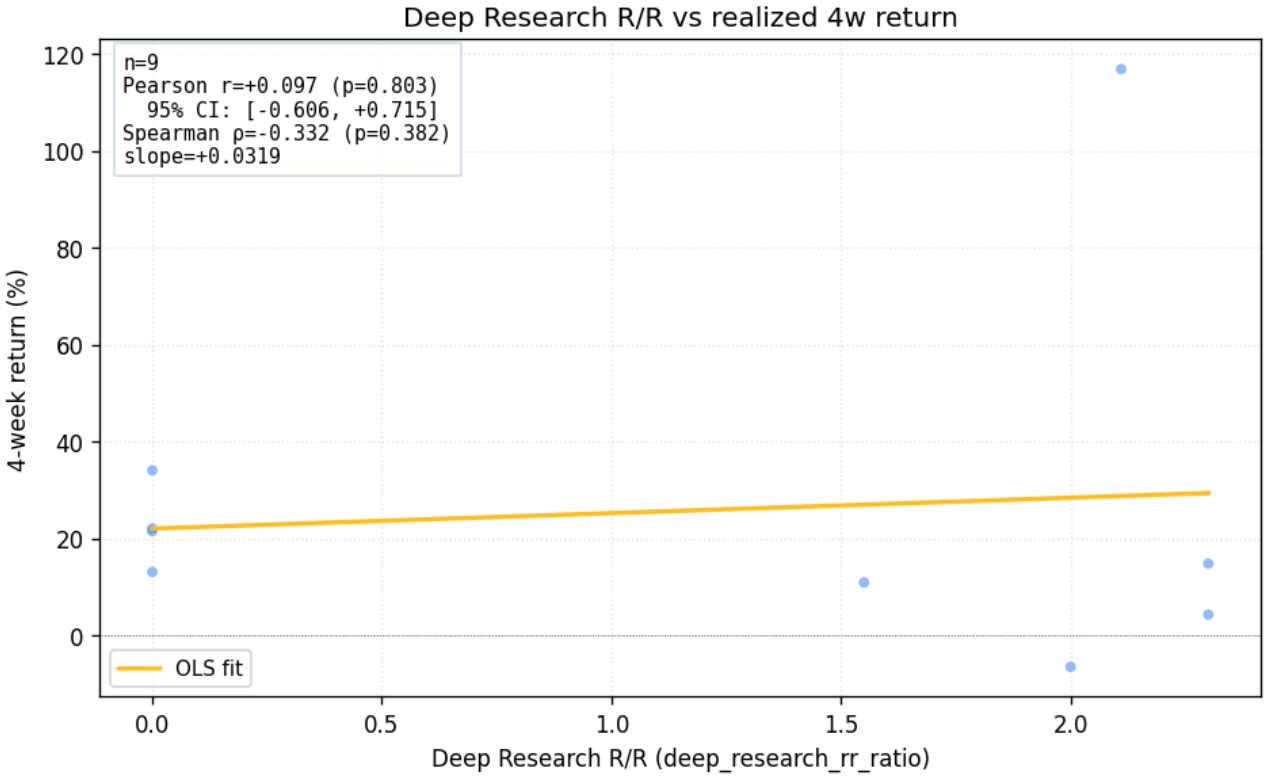

corr_dr

Interpretation. Pearson r captures linear association — if a few high-R/R, high-return rows dominate, Pearson can be inflated even when most of the data is uncorrelated. Spearman ρ ranks the values and is robust to those outliers; if the two coefficients disagree sharply the relationship is not monotonic and shouldn't be treated as predictive.

5. Recovery patterns

- **AVOID** (n=188): 49% recovered, median 2 days; +20d post-recovery: +33.18%
- **NEUTRAL** (n=72): 57% recovered, median 4 days; +20d post-recovery: +4.66%
- **ENTER_LIMIT** (n=50): 34% recovered, median 1 days; +20d post-recovery: +15.83%
- **ENTER_NOW** (n=49): 47% recovered, median 1 days; +20d post-recovery: +16.59%

5.1 By AI council intent

group	n_total	n_recov	recov%	p25 d	p50 d	p75 d	p90 d	post +5d	post +10d	post +20d
AVOID	188	93	49%	0	2	4	8	+1.04%	-0.57%	+33.18%
NEUTRAL	72	41	57%	1	4	8	10	-0.58%	+0.25%	+4.66%
ENTER_LIMIT	50	17	34%	0	1	5	10.8	+2.25%	+2.08%	+15.83%
ENTER_NOW	49	23	47%	0	1	4.5	7.8	+3.31%	+2.02%	+16.59%

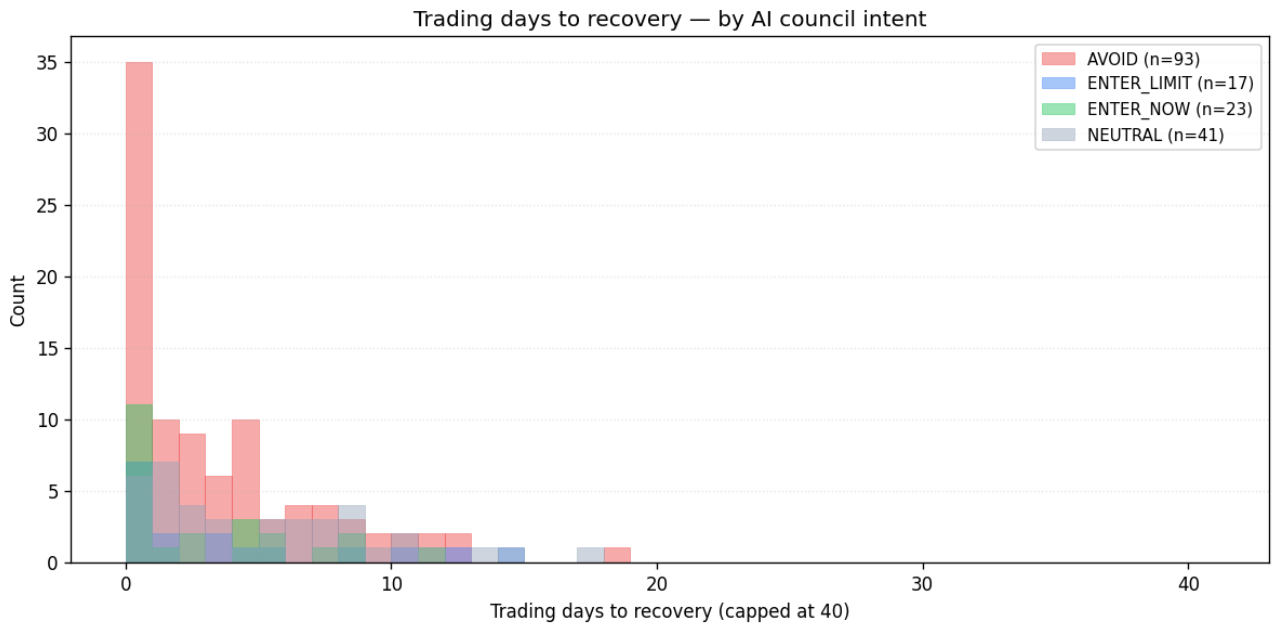

recovery_hist

5.2 By Deep Research verdict

group	n_total	n_recov	recov%	p25 d	p50 d	p75 d	p90 d	post +5d	post +10d	post +20d
nan	291	146	50%	0	2	6	9.5	+0.55%	-0.30%	+26.12%
BUY_LIMIT	32	13	41%	0	4	5	10.4	+5.13%	+1.90%	+15.58%
AVOID	26	9	35%	0	0	4	7.2	+5.56%	+4.42%	+15.45%
BUY	9	5	56%	0	2	3	6	+2.38%	+2.97%	
WATCH	1	1	100%	0	0	0	0	-0.41%		

Interpretation. `days_to_recover` is the number of trading days from the decision date until the price first reaches the pre-drop level. The post-recovery columns measure what the stock did over the next 5/10/20 trading days *after* recovery — a positive number means the stock kept going up after reaching its pre-drop level.

6. Performance over time vs S&P 500

Median return path from the decision date forward, with SPY's median over the same calendar windows (dashed line) as a passive benchmark.

6.1 By AI council intent

Two views: median path with **inter-quartile range** band (default), and mean path with **95% t-CI** band.

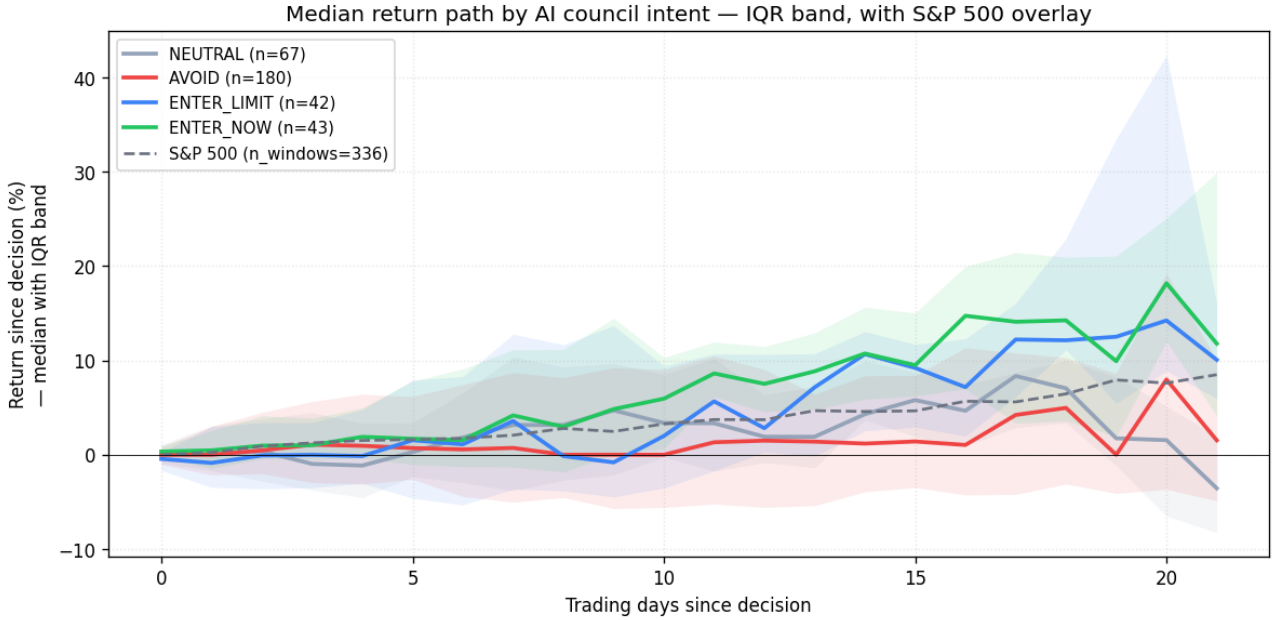

ts_intent

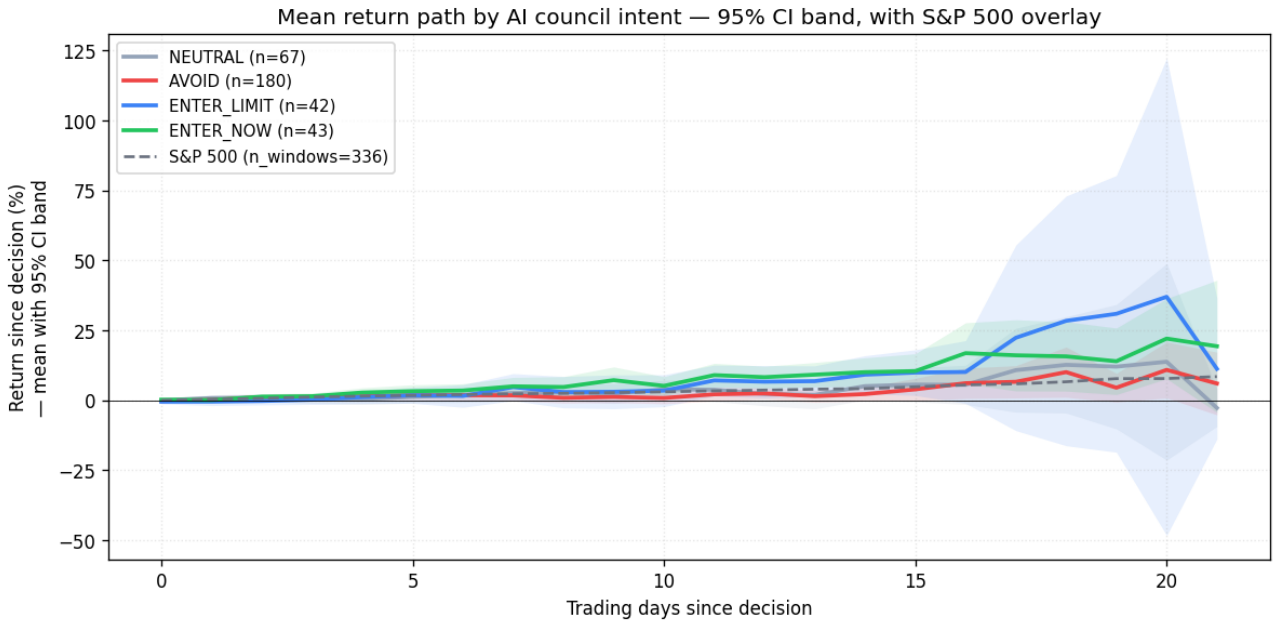

ts_intent_mean

6.2 By Deep Research verdict

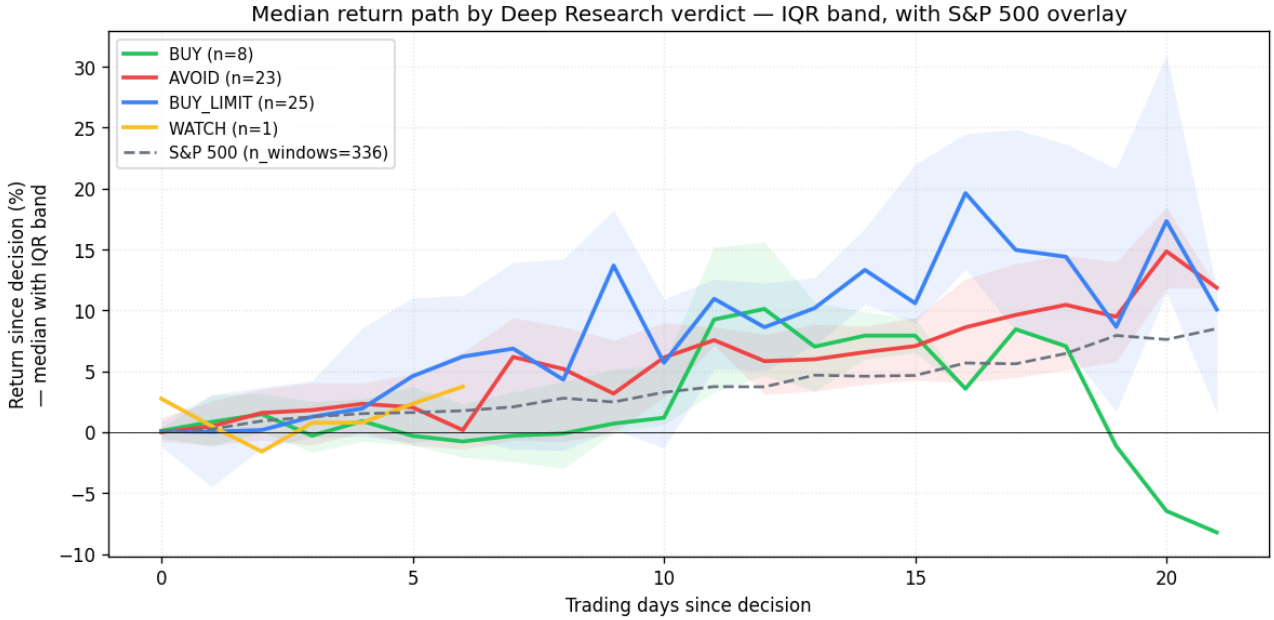

ts_dr

6.3 Excess return vs SPY (alpha)

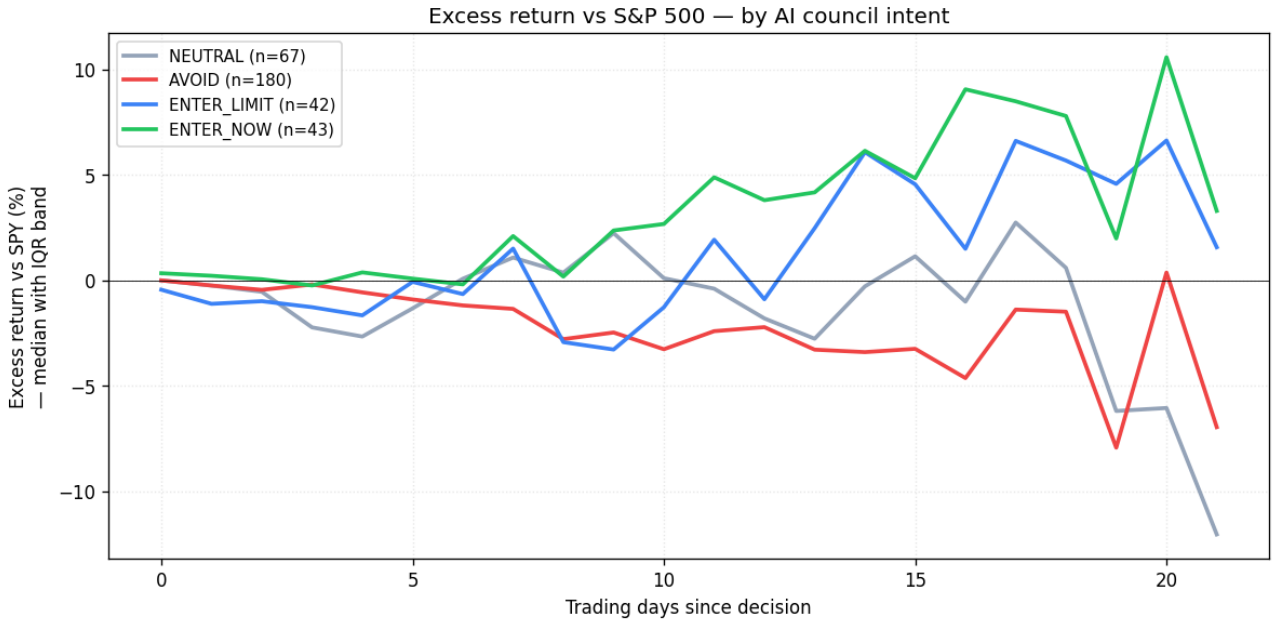

ts_alpha

6.4 Per-decision BUY trajectories

Light grey lines are individual ENTER_NOW + ENTER_LIMIT decisions; bold lines are the per-intent medians. Useful for sense-checking how typical the median trajectory really is.

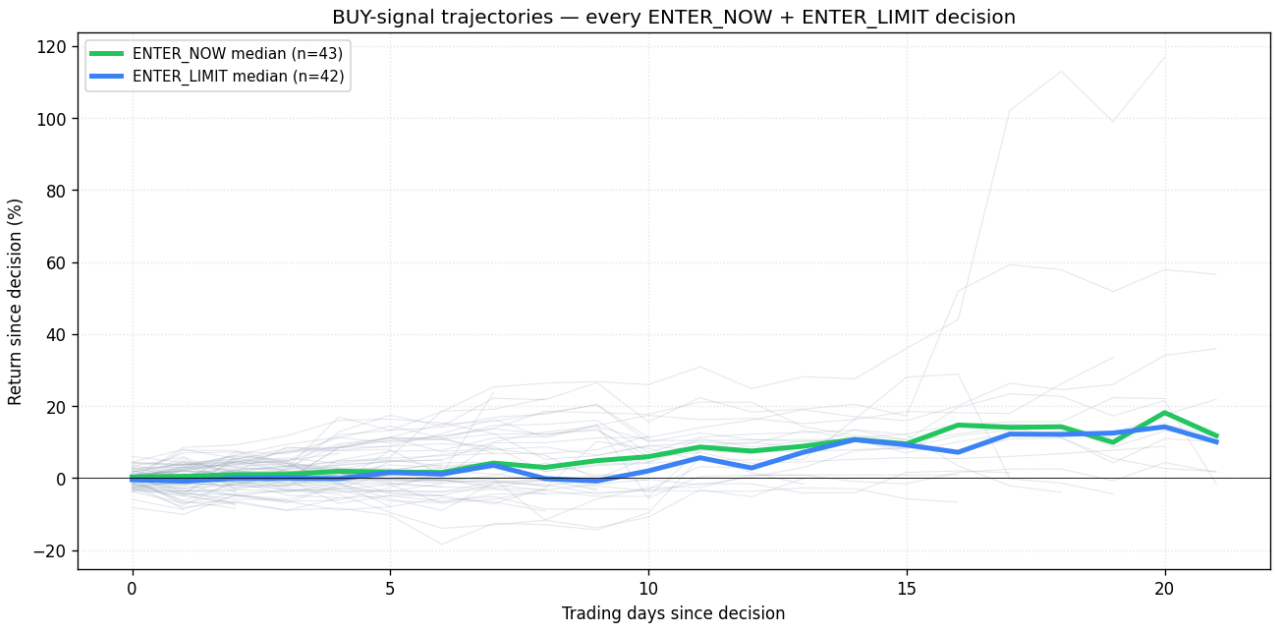

ts_spaghetti

7. Drop-size buckets

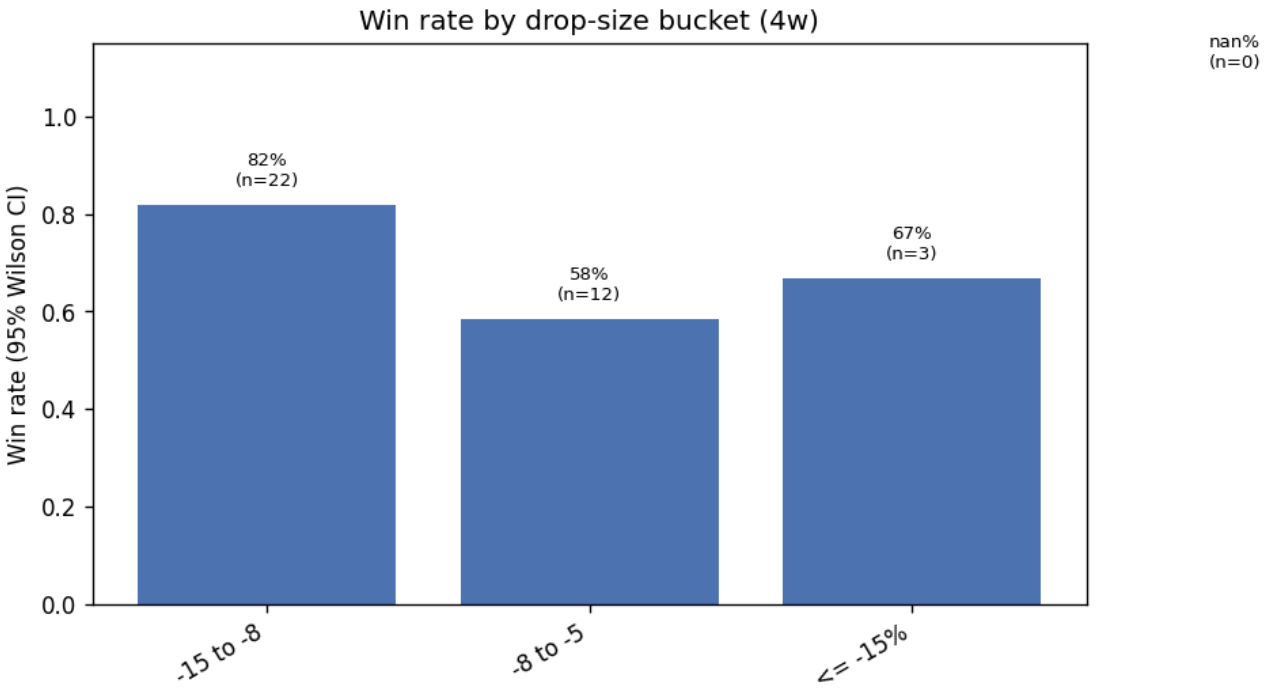

wr_drop

7a. Horizon comparison — 1w / 2w / 4w

The 4-week return is small-sample (n=39 with completed bars). 1-week and 2-week returns are much better powered (n=272 and n=184). All subsequent significance tests and bucket aggregations are now computed at every horizon; the bars below show win rate (top) and average return (bottom) side-by-side per intent. Wilson + t-CI error bars come along for the ride.

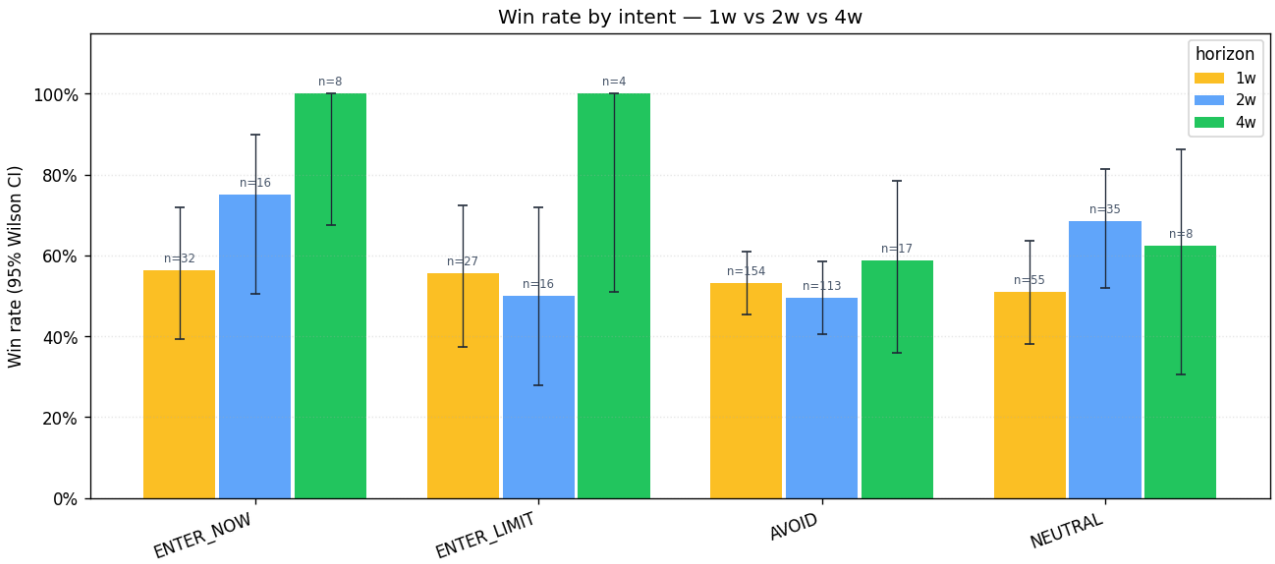

wr_intent_multi

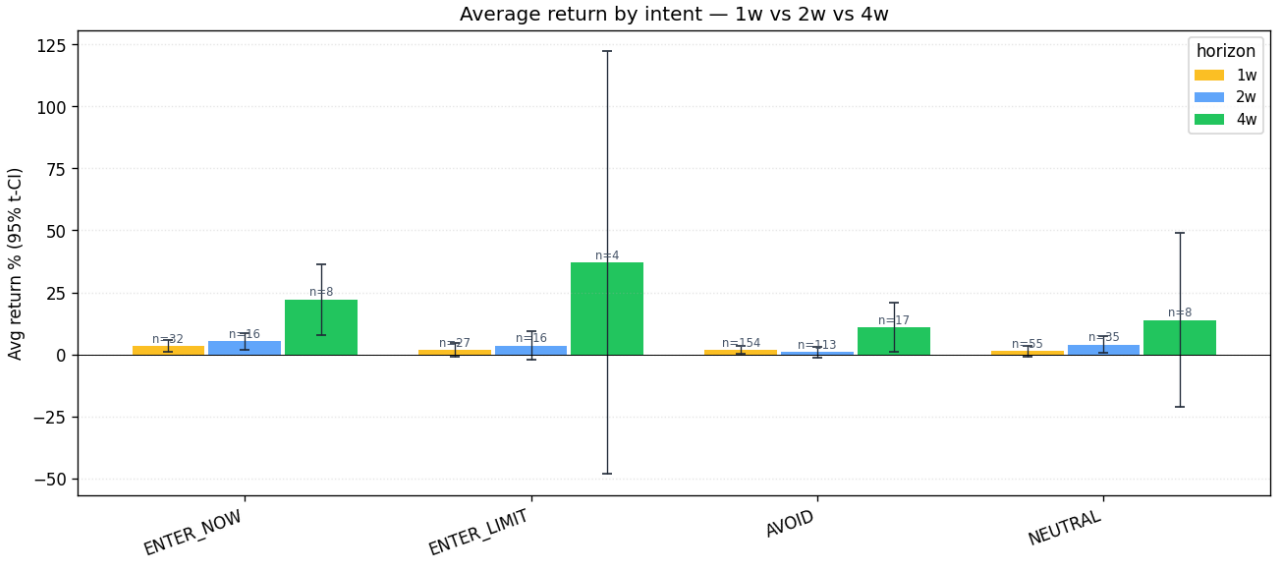

ar_intent_multi

Per-intent win rate and avg return at each horizon:

group	horizon	n	win_rate	WR 95% CI	avg_return	AR 95% CI
AVOID	1w	154	53.2%	[45.4%, 61.0%]	+1.82%	[+0.27%, +3.36%]
AVOID	2w	113	49.6%	[40.5%, 58.6%]	+0.85%	[-1.20%, +2.90%]
AVOID	4w	17	58.8%	[36.0%, 78.4%]	+10.89%	[+0.89%, +20.89%]
ENTER_LIMIT	1w	27	55.6%	[37.3%, 72.4%]	+1.82%	[-1.03%, +4.67%]
ENTER_LIMIT	2w	16	50.0%	[28.0%, 72.0%]	+3.46%	[-2.25%, +9.17%]
ENTER_LIMIT	4w	4	100.0%	[51.0%, 100.0%]	+37.01%	[-48.27%, +122.30%]

group	horizon	n	win_rate	WR 95% CI	avg_return	AR 95% CI
ENTER_NOW	1w	32	56.2%	[39.3%, 71.8%]	+3.28%	[+0.98%, +5.58%]
ENTER_NOW	2w	16	75.0%	[50.5%, 89.8%]	+5.18%	[+1.87%, +8.49%]
ENTER_NOW	4w	8	100.0%	[67.6%, 100.0%]	+22.08%	[+7.77%, +36.39%]
NEUTRAL	1w	55	50.9%	[38.1%, 63.6%]	+1.24%	[-0.95%, +3.43%]
NEUTRAL	2w	35	68.6%	[52.0%, 81.4%]	+3.94%	[+0.71%, +7.17%]
NEUTRAL	4w	8	62.5%	[30.6%, 86.3%]	+13.80%	[-21.33%, +48.93%]

By drop-size bucket at each horizon:

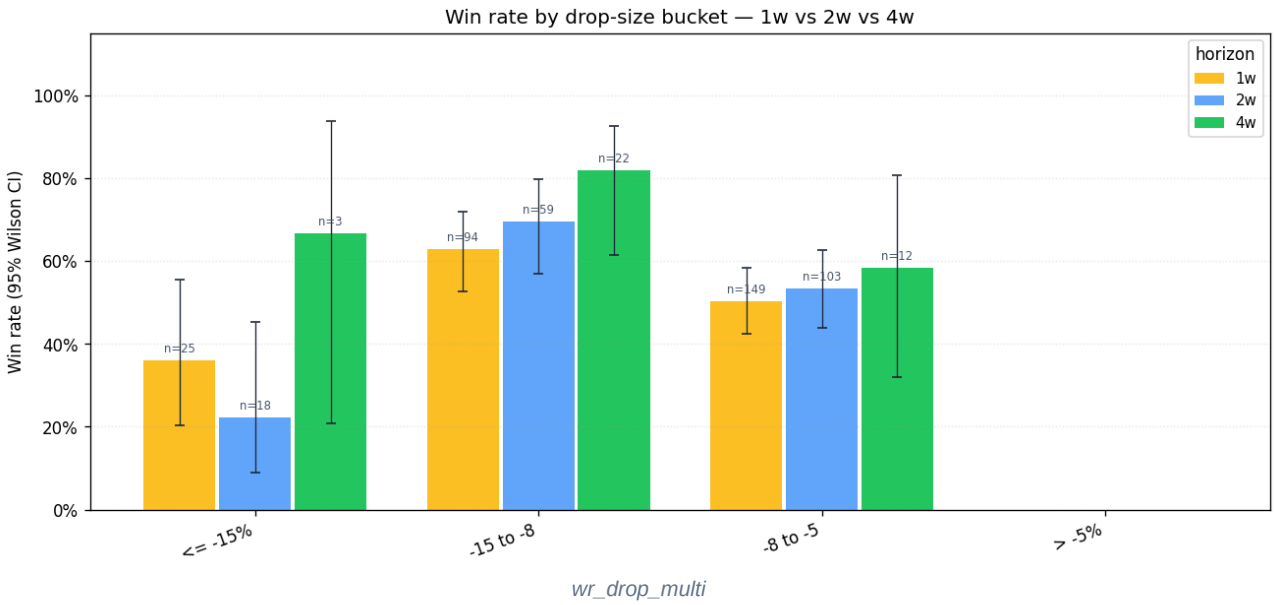

By PM R/R bucket at each horizon:

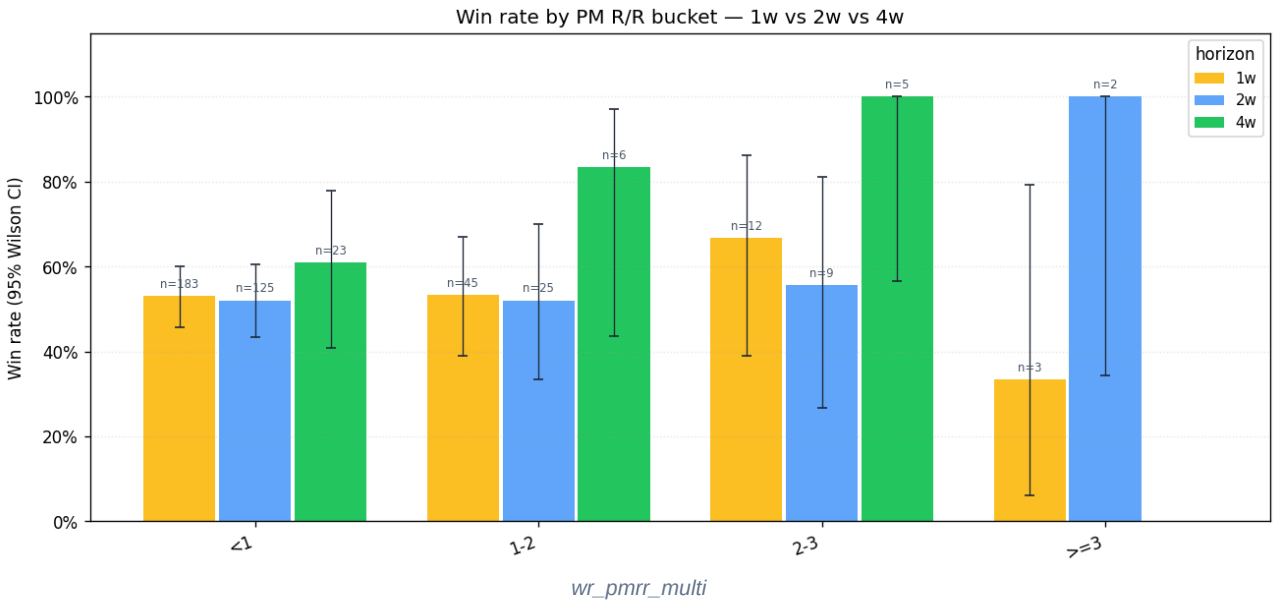

By DR R/R bucket at each horizon:

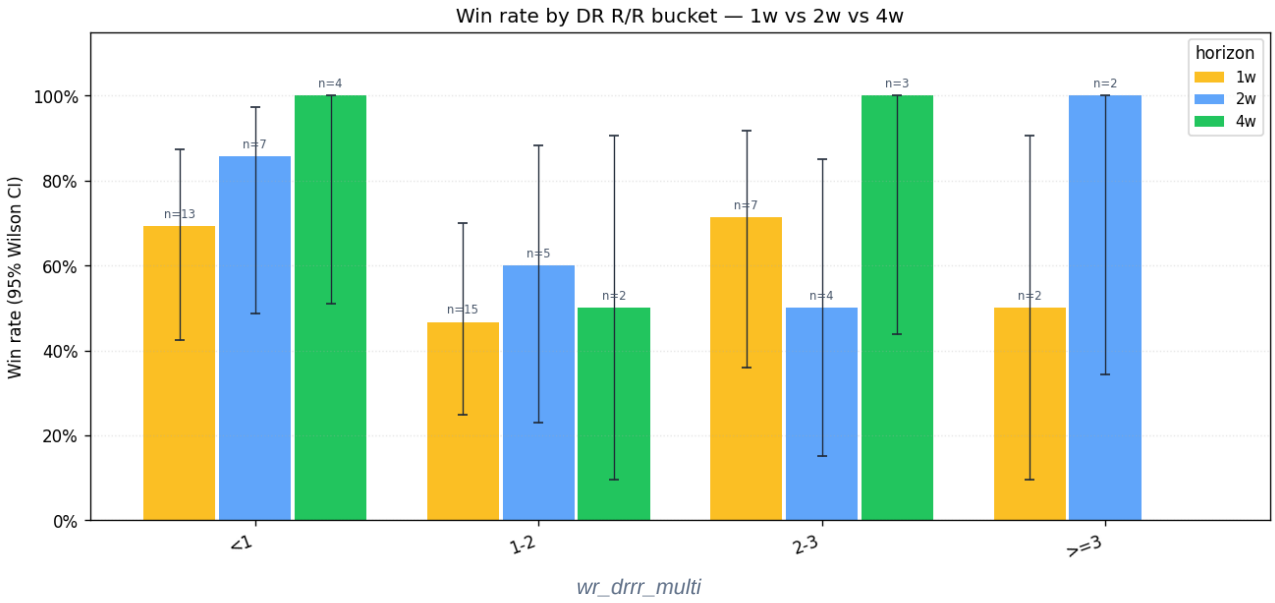

8a. R/R distribution by verdict (categorical correlation)

How does each council assign its R/R ratings to its own verdict groups? Below: per-group descriptives, one-way ANOVA (parametric), and Kruskal-Wallis (rank-based) for the omnibus test of "are the group distributions the same?" Plus pairwise Welch t-tests with FDR-adjusted p-values.

8a.1 PM R/R by AI council intent

group	n	mean	SE	CI low	CI high	median	std	min	max
AVOID	176	0.643	0.044	0.556	0.731	0.5	0.589	0	5.4
NEUTRAL	51	0.886	0.084	0.718	1.054	0.7	0.598	0.1	2.8
ENTER_LIMIT	50	1.144	0.095	0.952	1.335	1	0.673	0	2.9
ENTER_NOW	49	1.503	0.127	1.249	1.758	1.56	0.887	0	4

Omnibus: ANOVA: F=25.09, p=<0.001 · Kruskal-Wallis: H=65.20, p=<0.001.

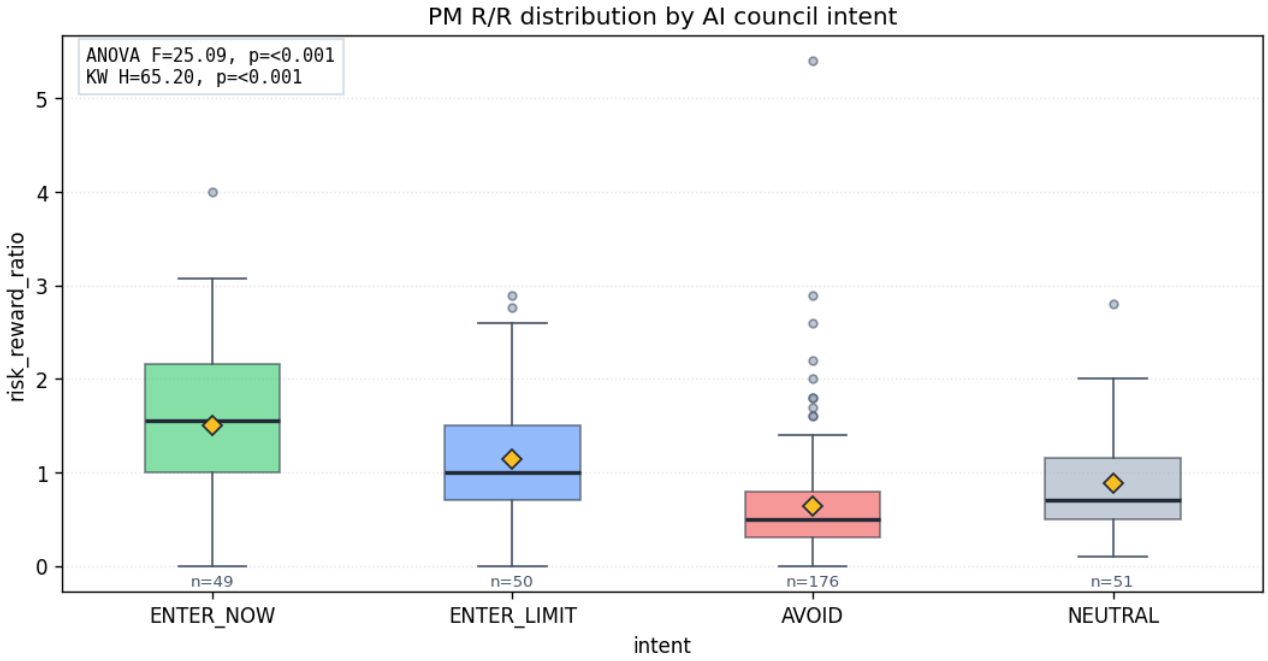

rr_box_pm

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	ENTER_LIMIT	176	50	-50.04%	-0.82	<0.001	<0.001	<0.001	<0.001	■
AVOID	ENTER_NOW	176	49	-86.03%	-1.29	<0.001	<0.001	<0.001	<0.001	■
AVOID	NEUTRAL	176	51	-24.31%	-0.41	0.012	0.018	0.002	0.003	■
ENTER_LIMIT	ENTER_NOW	50	49	-35.99%	-0.46	0.026	0.031	0.010	0.013	■
ENTER_LIMIT	NEUTRAL	50	51	+25.73%	0.4	0.045	0.045	0.022	0.022	■
ENTER_NOW	NEUTRAL	49	51	+61.72%	0.82	<0.001	<0.001	<0.001	<0.001	■

8a.2 DR R/R by Deep Research verdict

group	n	mean	SE	CI low	CI high	median	std	min	max
BUY_LIMIT	32	1.634	0.131	1.366	1.902	1.675	0.743	0	3.07
AVOID	23	1.231	0.184	0.849	1.614	1.03	0.884	0	2.9
BUY	8	1.952	0.402	1.002	2.903	2	1.137	0	4
WATCH	1	0.82				0.82		0.82	0.82

Omnibus: ANOVA: F=2.65, p=0.079 · Kruskal-Wallis: H=5.28, p=0.071.

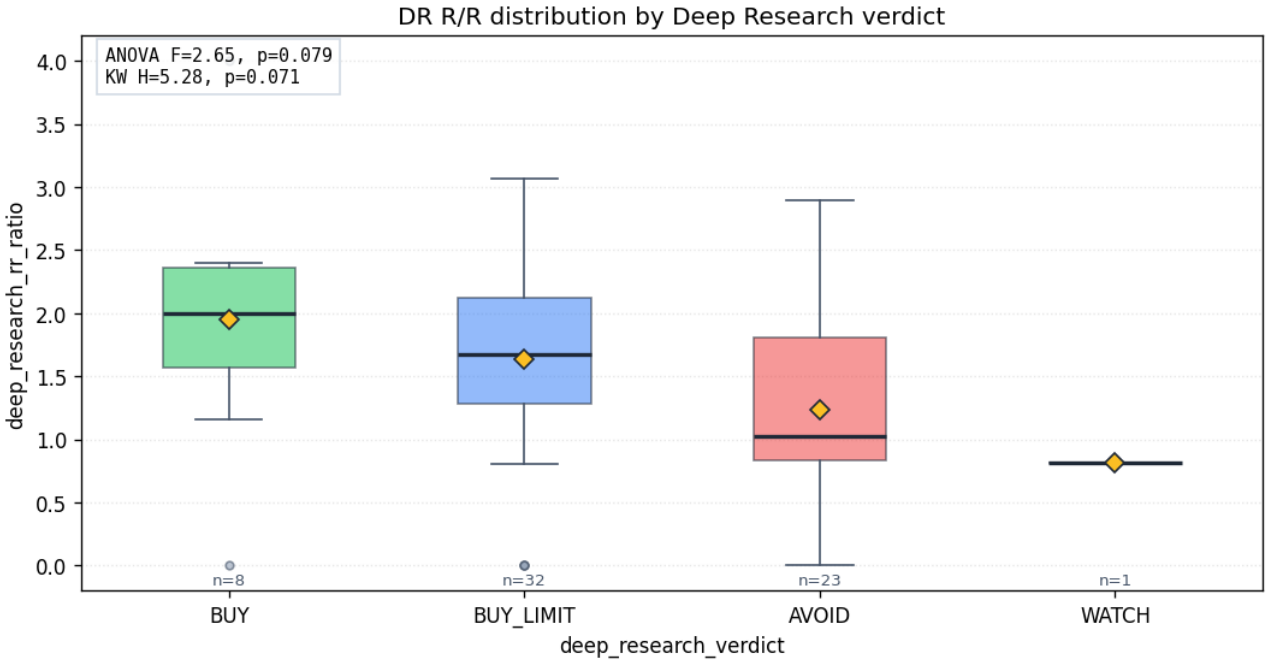

rr_box_dr

A	B	n_A	n_B	Δ mean	Cohen d	Welch p	Welch p (FDR)	MWU p	MWU p (FDR)	Sig?
AVOID	BUY	23	8	-72.12%	-0.76	0.134	0.201	0.07	0.105	—
AVOID	BUY_LIMIT	23	32	-40.28%	-0.5	0.082	0.201	0.061	0.105	—
BUY	BUY_LIMIT	8	32	+31.84%	0.38	0.472	0.472	0.37	0.37	—

Interpretation. ANOVA and Kruskal-Wallis answer the question "is *any* group different from the others?" If both omnibus tests fail to reach p<0.05 the council is not assigning systematically different R/R to different verdict groups (and the pairwise tests will reflect that).

8b. High-R/R decisions in the cohort

Top 25 rows by R/R ratio for each council, with their realized returns where available. Full lists exported as `data/top_pm_rr.csv` and `data/top_dr_rr.csv`.

8b.1 Top by PM R/R (`risk_reward_ratio`)

symbol	decision_date	intent	recommendation	risk_reward_ratio	drop_percent	price_at_decision	deep_research_verdict	return_1w	return_2w	return_4w	max_roi_4w	max_drawdown_4w
CHKP	2026-04-30	AVOID	AVOID	5.4	-1694.06%	$116.25		-0.59%			+7.24%	-3.46%
LBRDP	2026-04-24	ENTER_NOW	BUY	4	-597.17%	$21.96	BUY	-0.30%	+0.89%		+6.58%	-4.17%

symbol	decision_date	intent	recommendation	risk_reward_ratio	drop_percent	price_at_decision	deep_research_verdict	return_1w	return_2w	return_4w	max_roi_4w	max_drawdown_4w
NICE	2026-04-11	ENTER_NOW	BUY	3.07	-714.15%	$97.00	BUY_LIMIT	+8.79%	+4.08%		+30.41%	-6.49%
NOC	2026-04-21	AVOID	AVOID	2.9	-580.15%	$618.87		-6.63%	-9.74%		+5.52%	-12.06%
FDX	2026-05-04	ENTER_LIMIT	BUY_LIMIT	2.9	-674.68%	$367.11	AVOID				+4.41%	-3.53%
TEAM	2026-04-09	ENTER_NOW	BUY	2.8	-804.78%	$58.50		+17.49%	+15.59%	+57.90%	+64.65%	-4.26%
TRUMF	2026-04-13	NEUTRAL	WATCH	2.8	-528.93%	$12.40		+1.77%	+1.77%		+8.47%	-0.00%
CRC	2026-05-06	ENTER_LIMIT	BUY_LIMIT	2.76	-960.36%	$63.40	AVOID				+8.76%	-7.80%
CE	2026-05-06	ENTER_LIMIT	BUY_LIMIT	2.6	-954.93%	$62.42	BUY_LIMIT				+6.92%	-10.29%
FIS	2026-05-08	ENTER_NOW	BUY	2.6	-584.13%	$44.49	BUY_LIMIT					
MHK	2026-05-04	AVOID	AVOID	2.6	-546.88%	$94.47					+13.50%	-0.92%
ALC	2026-05-06	ENTER_NOW	BUY	2.4	-967.09%	$67.25	BUY				+0.61%	-7.21%
HUBS	2026-05-08	ENTER_LIMIT	BUY_LIMIT	2.4	-2333.01%	$186.86	BUY_LIMIT					
REGN	2026-04-29	ENTER_NOW	BUY	2.35	-621.85%	$686.26	BUY	+5.07%			+5.81%	-2.54%
IAG	2026-04-21	ENTER_NOW	BUY	2.33	-718.53%	$17.18	BUY_LIMIT	-4.02%	-5.06%		+13.10%	-6.69%

symbol	decision_date	intent	recommendation	risk_reward_ratio	drop_percent	price_at_decision	deep_research_verdict	return_1w	return_2w	return_4w	max_roi_4w	max_drawdown_4w
POOL	2026-05-05	ENTER_LIMIT	BUY_LIMIT	2.3	-606.21%	$190.91	AVOID				+3.88%	-3.36%
ADSK	2026-04-09	ENTER_NOW	BUY	2.3	-916.27%	$218.60	AVOID	+11.24%	+6.12%	+14.84%	+16.72%	-2.06%
NOW	2026-04-09	ENTER_NOW	BUY	2.3	-793.58%	$89.73	BUY_LIMIT	+7.47%	-5.52%	+4.30%	+16.45%	-9.47%
WTW	2026-04-30	ENTER_NOW	BUY	2.24	-1417.39%	$248.99	BUY_LIMIT	+3.65%			+7.47%	-0.96%
CEBCF	2026-04-10	AVOID	AVOID	2.2	-1436.36%	$0.38		+8.81%	+16.77%	+8.81%	+16.77%	+0.00%
CHYM	2026-05-07	ENTER_NOW	BUY	2.2	-704.10%	$20.20					+6.19%	-9.90%
PPERF	2026-04-11	ENTER_NOW	BUY	2.2	-2316.78%	$0.26	AVOID	-0.00%	-0.00%		-0.00%	-0.00%
ZBH	2026-04-28	ENTER_NOW	BUY	2.16	-706.34%	$86.05	BUY_LIMIT	-3.53%			+3.01%	-7.23%
DOCN	2026-04-10	ENTER_LIMIT	BUY_LIMIT	2.11	-1328.44%	$75.59	BUY_LIMIT	+13.28%	+25.96%	+116.89%	+117.98%	-4.52%
TOST	2026-05-08	ENTER_NOW	BUY	2.11	-1463.58%	$25.08	BUY_LIMIT					

8b.2 Top by DR R/R (`deep_research_rr_ratio`)

symbol	decision_date	intent	recommendation	deep_research_rr_ratio	drop_percent	price_at_decision	deep_research_verdict	return_1w	return_2w	return_4w	max_roi_4w	max_drawdown_4w
LBRDP	2026-04-24	ENTER_NOW	BUY	4	-597.17%	$21.96	BUY	-0.30%	+0.89%		+6.58%	-4.17%
NICE	2026-04-11	ENTER_NOW	BUY	3.07	-714.15%	$97.00	BUY_LIMIT	+8.79%	+4.08%		+30.41%	-6.49%
FDX	2026-05-04	ENTER_LIMIT	BUY_LIMIT	2.9	-674.68%	$367.11	AVOID				+4.41%	-3.53%
CRC	2026-05-06	ENTER_LIMIT	BUY_LIMIT	2.76	-960.36%	$63.40	AVOID				+8.76%	-7.80%
FIS	2026-05-08	ENTER_NOW	BUY	2.6	-584.13%	$44.49	BUY_LIMIT					
CE	2026-05-06	ENTER_LIMIT	BUY_LIMIT	2.6	-954.93%	$62.42	BUY_LIMIT				+6.92%	-10.29%
HUBS	2026-05-08	ENTER_LIMIT	BUY_LIMIT	2.4	-2333.01%	$186.86	BUY_LIMIT					
ALC	2026-05-06	ENTER_NOW	BUY	2.4	-967.09%	$67.25	BUY				+0.61%	-7.21%
REGN	2026-04-29	ENTER_NOW	BUY	2.35	-621.85%	$686.26	BUY	+5.07%			+5.81%	-2.54%
IAG	2026-04-21	ENTER_NOW	BUY	2.33	-718.53%	$17.18	BUY_LIMIT	-4.02%	-5.06%		+13.10%	-6.69%
NOW	2026-04-09	ENTER_NOW	BUY	2.3	-793.58%	$89.73	BUY_LIMIT	+7.47%	-5.52%	+4.30%	+16.45%	-9.47%
POOL	2026-05-05	ENTER_LIMIT	BUY_LIMIT	2.3	-606.21%	$190.91	AVOID				+3.88%	-3.36%
ADSK	2026-04-09	ENTER_NOW	BUY	2.3	-916.27%	$218.60	AVOID	+11.24%	+6.12%	+14.84%	+16.72%	-2.06%

symbol	decision_date	intent	recommendation	deep_research_rr_ratio	drop_percent	price_at_decision	deep_research_verdict	return_1w	return_2w	return_4w	max_roi_4w	max_drawdown_4w
WTW	2026-04-30	ENTER_NOW	BUY	2.24	-1417.39%	$248.99	BUY_LIMIT	+3.65%			+7.47%	-0.96%
ZBH	2026-04-28	ENTER_NOW	BUY	2.16	-706.34%	$86.05	BUY_LIMIT	-3.53%			+3.01%	-7.23%
TOST	2026-05-08	ENTER_NOW	BUY	2.11	-1463.58%	$25.08	BUY_LIMIT					
DOCN	2026-04-10	ENTER_LIMIT	BUY_LIMIT	2.11	-1328.44%	$75.59	BUY_LIMIT	+13.28%	+25.96%	+116.89%	+117.98%	-4.52%
SYM	2026-05-08	ENTER_NOW	BUY	2.03	-845.30%	$51.66	BUY_LIMIT					
ALLE	2026-04-28	ENTER_NOW	BUY	2	-745.96%	$137.33	BUY_LIMIT	-3.67%			+2.93%	-4.45%
APA	2026-04-09	NEUTRAL	PENDING	2	-979.98%	$38.75	BUY	-2.19%	-0.23%	-6.48%	+8.62%	-13.73%
NMR	2026-04-24	ENTER_NOW	BUY	2	-526.00%	$7.83	BUY	-0.32%	+1.47%		+4.79%	-1.47%
DPZ	2026-04-27	ENTER_NOW	BUY	1.96	-934.67%	$333.45	BUY_LIMIT	-0.91%			+3.53%	-3.65%
EMBJ	2026-05-08	ENTER_LIMIT	BUY_LIMIT	1.9	-791.14%	$62.39	AVOID					
CLX	2026-05-01	ENTER_LIMIT	BUY_LIMIT	1.86	-895.38%	$87.81	AVOID	+4.96%			+6.44%	-3.54%
EXPE	2026-05-08	ENTER_NOW	BUY	1.8	-869.50%	$230.81	BUY_LIMIT					

9. Profit and loss decomposition

Two complementary views of how each group's wins and losses played out *over time*.

8.1 Winner vs loser trajectories per category

Each cohort row is classified at day-20 by the sign of its return. The panels below show the **mean winning trajectory** vs the **mean losing trajectory** within each category, with 95% CI bands.

Why this matters: a high overall avg return can come from many small wins or a few large ones; this chart lets you see the asymmetry.

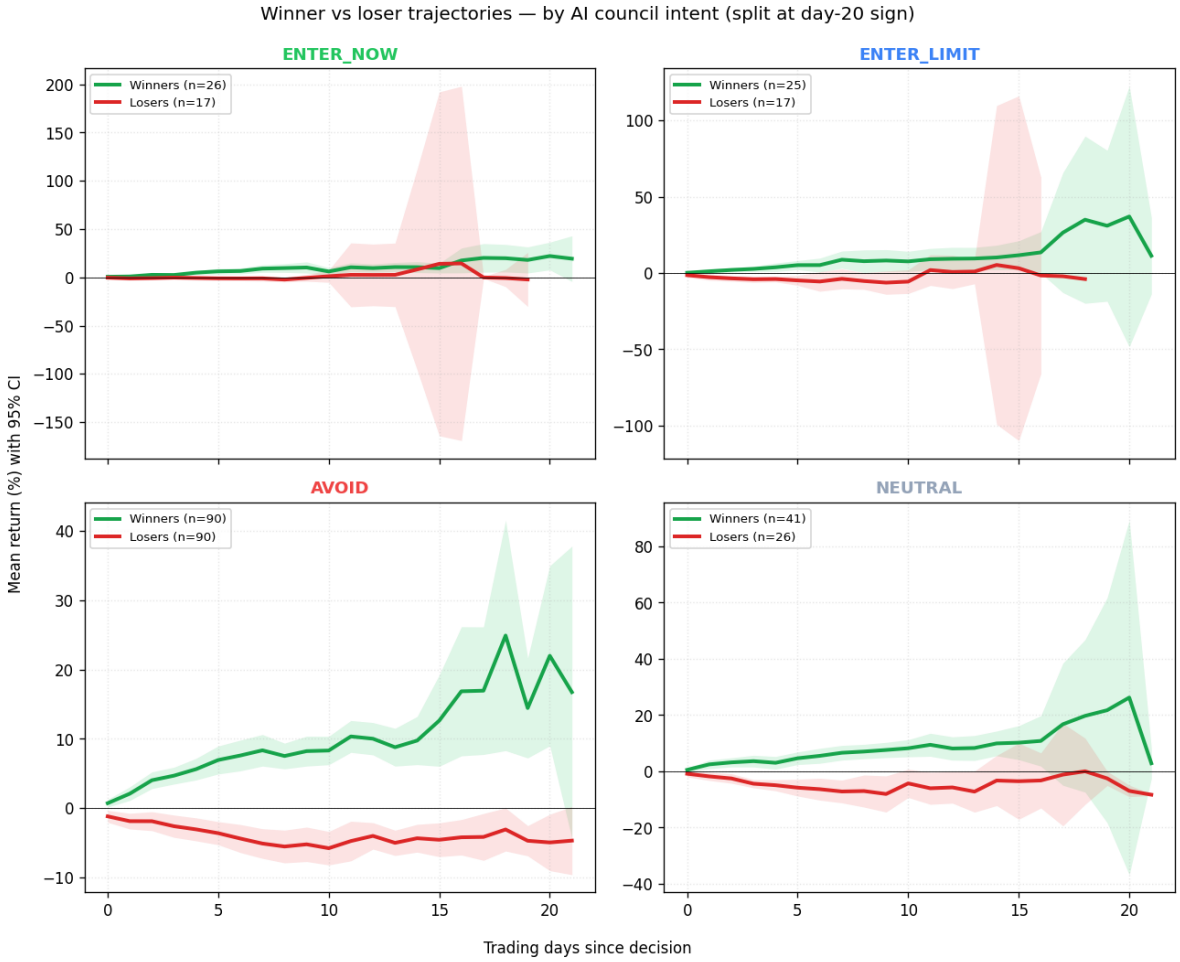

winloss_intent

winloss_dr

8.2 Cumulative mark-to-market P&L over calendar time

Assume **\$1 is invested at every signal** at the decision-date close, held forward, and marked to its closing price every subsequent trading day. The lines below sum that mark-to-market P&L across every open position, by category, on each calendar date.

Useful for seeing *when* the P&L accrued (early jump? steady drift?) and how each category's book performed in real time.

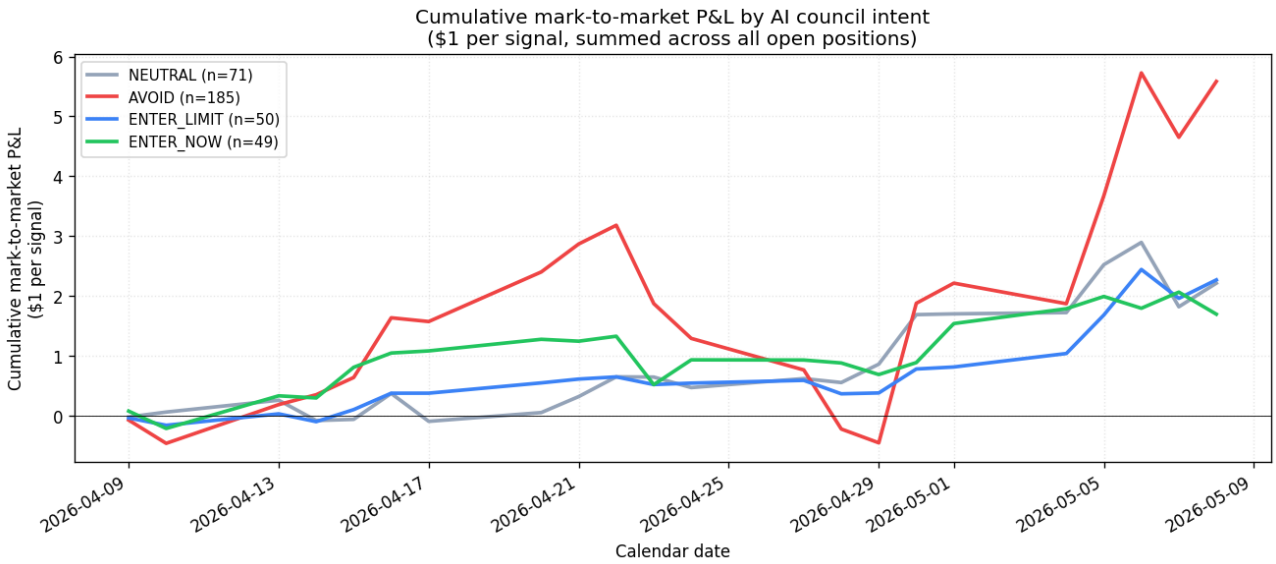

cum_pnl_intent

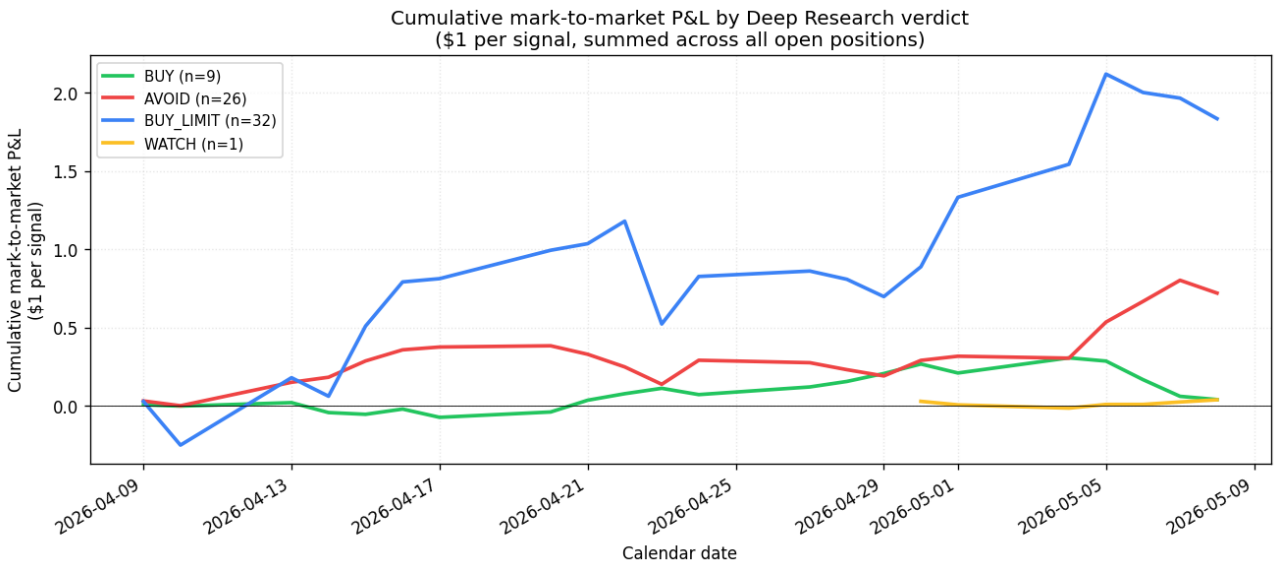

cum_pnl_dr

10. Limitations

- **Forward-window coverage.** With current `decision_date` range, no decision has more than ~22 trading days of forward data, which means the 4-week and 8-week return columns are NaN for most rows. Re-running this script after more time elapses extends every horizon naturally.
- **Sample size.** After dropping rows without 4w returns, intent groups have n=4–18 and DR-verdict groups have n=1–6. The pairwise significance tests are honest about this — they refuse to call differences "real" until the data catches up.
- **Market regime.** Cohort window appears to coincide with a broad SPY rally (+7.6% median over 20 trading days). Many AVOIDs would have been profitable passive holdings; that is a property of this regime and should not be generalized.
- **Storage duplication.** `deep_research_action` and `deep_research_verdict` carry identical values in this DB; the Q2/3.1 sections are therefore redundant against the underlying signal.

11. Recommendations

- **Wait, then re-run.** The single largest analytical lift is more time. Once the earliest decisions reach their 8-week mark, re-run `build_package.py` and the same charts will tell a much sharper story.
- **Investigate AVOID hits.** AVOIDs with high `+20d` post-recovery returns are worth pulling individually — was the AVOID a calibration bug, or did the model correctly price in higher risk that didn't materialize this regime?
- **Drop the duplicate column.** Either consolidate `deep_research_verdict` and `deep_research_action`, or document why both exist.
- **Backfill `ai_score`.** Currently populated for 10/363 rows, all = 50. Either fully populate or remove from prompts; right now it can't inform any analysis.

Appendix

All raw data underlying this report is in `data/`:

- `cohort_enriched.csv` — every decision with computed return columns
- `winrate_by_*.csv` — per-group aggregations
- `stats_*.csv|.json` — significance and correlation results
- `time_series.json` — per-day median paths
- `time_series_individuals.json` — every BUY-signal trajectory
- `spy_bars.csv` — SPY OHLC for the cohort window
- `full_payload.json` — the entire JSON payload that drives `deep-dive.html`

Interactive HTML report: `deep-dive.html`